

ASSIGNMENT 2 OF 2 · REQUIRED COVER SHEET · DUE SUNDAY 13 SEPTEMBER 2026, 23:59 SGT

A2 Team Self-Appraisal

Team ID: B-1 Section: B Problem chosen: A

Members: Meng Sijia; Shi Shuyi; Su Yang; Isha Kirti Ghia; Sun Hanyu; Zhang Jiayang

Repository URL: https://github.com/leos5723-ys/NTU_6201_Group-1 Date: 2026.09.20

ONE SHEET PER TEAM, not one per member. Complete it together, in one sitting, after the work is done — it takes about twenty minutes. CHECK IT INTO YOUR A2 SUBMISSION FOLDER alongside the rest of the submission, by the A2 deadline of Sunday 13 September 2026, 23:59 SGT; do not email it. This sheet is required and is not graded, but a missing self-appraisal is treated as an incomplete submission.

1 · Rate your team against Rubric 1

Circle or mark one band per criterion. Be honest rather than generous — a team that marks itself Excellent everywhere tells the marker nothing, and a large gap between your rating and the marker's is the single most useful thing on this page.

Criterion	Weight	Excellent (4)	Proficient (3)	Developing (2)	Limited (1)
Conceptual Understanding	25%	4			
Technical Execution	30%	4			
Reasoning & Justification	25%		3		
Communication & Clarity	20%	4			

2 · The one decision you would defend hardest

Name a single design decision — a tool you cut, the autonomy setting you chose, the dependency rule behind your parallel calls, the model you settled on — and say what evidence made it. Four or five sentences.

We would defend the dependency rule behind our parallel calls. The starter example paired a policy lookup with coverage calls that needed the policy ID it had not yet returned. Our line-review interface instead uses the member ID already available from the claim and resolves the policy internally. The scripted comparison retained 40/40 correct cases while reducing turns from 251 to 157 and estimated input tokens from 911,506 to 609,144. This is measured efficiency on the scripted trajectories, not proof that every live model will choose the same ordering.

3 · The one you are least sure about

Every team has one. Naming it costs you nothing here and it is often the most interesting thing in the submission.

Our hostile-input protection is deliberately narrow. Known instruction patterns are detected in code before the claim observation returns to the model. Passing these cases demonstrates the shipped guard, not independent model resistance to prompt injection. We would require broader adversarial testing and independent security review before deployment, including paraphrases that do not match the current patterns.

4 · Your headline numbers

Copy these from your repository. If any cell is empty, that is worth noticing before you submit rather than after.

	Value		Value
Evaluation cases	40	Guardrail cases	10 scripted checklist cases
Negative cases	10	Live runs per model (D4)	60: ordinary ×1; negative ×3
Pass rate (best model)	DeepSeek v2: 55/60 (91.7%)	Pass rate on negatives only	DeepSeek v2: 29/30 (96.7%)
Models in the battery	5 v2 model families + Gemini v1 control	Median turns per run	DeepSeek: 4
The live model YOU ran (D5b)	See member/model table below	Your pass rate on it	See member/model table below
Turns saved by parallel calls	94 total turns: 251 → 157 on 40 scripted cases	Cost per successful task	\$0.82758 per incoming task resolved by agent or fallback; excludes fixed cost
Monthly cost at problem volume	\$7,020.66 at 8,000 claims/month; includes \$400 fixed cost	Break-even success rate	Qwen: 89.12% required; 47.50% measured
Tokens per call — v1 → v2	Gemini mean observation estimate: 82.37 → 63.42	Pass rate — v1 → v2, same model	Gemini: 33/60 → 20/60

Every member runs one live model, or the v1 pass — three models is the floor, N-1 for a team of N is the shape expected. Name YOURS. If you ran the v1 pass, write "v1 pass" and the model it sat on.

The last row is D2(b): the descriptor rewrite you measured. Two numbers each, before and after, on ONE model. An empty row here means the rewrite was never measured.

5 · Contribution

Who owned what. This should match CONTRIBUTIONS.md and the commit history. It is not used to scale marks — the team receives the team score — but it is one of the two signals the non-contribution clause relies on.

Member	Owned	Also contributed to
Meng Sijia	Case contracts; harness review; Gemini v2 battery	Live judgement; scripted CLM-8933; report review
Shi Shuyi	Case contracts; loop/tool review; Qwen v2 battery; integration	Evidence interpretation; calibration; report review
Su Yang	Case contracts; cost/ledger review; Claude v2 battery	Live judgement; scripted CLM-8952; report review
Isha Kirti Ghia	Case contracts; dependency review; Llama v2 battery	Repair proposals; live/scripted judgement; report review
Sun Hanyu	Case contracts; negative/hostile-input review; DeepSeek v2 battery	Live/scripted judgement; report review
Zhang Jiayang	Case contracts; ACI/tool review; Gemini v1 control	Live/scripted judgement; report corrections

Peer rating is separate from this sheet and comes AFTER it: NTU's Group Project peer-review function opens 13 September and closes Wednesday 16 September, 23:59 SGT. Each member completes it individually in NTULearn. It is a participation requirement.

6 · Declaration

Every member ticks every box. One sheet, signed by the team.

[☒] Every member of this team can explain every block of code we submit — what it does, and why it is there. (Block, not line: this is about function, not syntax.)

[☒] The pass rates, token counts, turn counts and costs we report are measurements we ran ourselves. None of them were supplied by an AI assistant.

[☒] Our submitted harness runs end to end on the scripted backend, with no key and no network.

[☒] We kept every fixture record we were given, unedited, and every case in our evaluation set carries a label. `check_my_data.py` passes on our data.

[☒] No part of this submission is drawn from any member's End-of-Course Project, and no part of it will be submitted as part of one.

[☒] This is our own work, and all sources, tools and assistance are attributed.

Signed on behalf of the team: Shi Shuyi **Date:** 2026.09.20

How this is used

This sheet is not marked and cannot lower your grade. It is read alongside your submission for two reasons. First, how a team reads its own work is usually the fastest route to feedback that is actually useful. Second, where a team's self-rating and the marking diverge sharply, that gap gets an explicit comment — which is where the most useful feedback tends to sit.

Member live-model results

Member	Selected model / version	Strict passes
Meng Sijia	Gemini 2.5 Flash Lite v2	20/60
Shi Shuyi	Qwen3 30B A3B Instruct v2	33/60
Su Yang	Claude Haiku 4.5 v2; complete rerun	35/60
Isha Kirti Ghia	Llama 4 Maverick v2	53/60
Sun Hanyu	DeepSeek V3.2 v2	55/60
Zhang Jiayang	Gemini 2.5 Flash Lite v1 pass	33/60